

Question 1

```
In [ ]: import numpy as np
```

```
In [ ]: def tanh(x):
    return np.tanh(x)

def tanh_grad(x):
    return 1 - np.tanh(x) ** 2

class NN():
    def __init__(self, architecture, learning_rate=0.1, activation=lambda x: x, activation_grad):
        '''This is a fully connected NN. The architecture is a list,
        with each element specifying the number of nodes in each layer'''
        self.arch = architecture
        self.num_layers = len(self.arch) - 1
        self.activation = activation
        self.activation_grad = activation_grad
        self.lr = learning_rate
        self.init_weights()

    def init_weights(self):
        np.random.seed(0) # don't modify this. GSI uses this for checking outputs
        self.weights = []
        self.biases = []
        for n in range(self.num_layers):
            self.weights.append(np.random.random((self.arch[n], self.arch[n+1])))
            self.biases.append(np.random.random((1, self.arch[n + 1])))

    def feed_forward(self, X):
        self.a_ns = []
        self.z_ns = []
        self.a_ns.append(X)
        for n in range(self.num_layers):
            z_n = np.dot(self.a_ns[-1], self.weights[n]) + self.biases[n]

            self.z_ns.append(z_n)
            self.a_ns.append(self.activation(z_n))

        return self.a_ns[-1]

    def loss_func(self, X, y):
        feed_forward = self.feed_forward(X)
        loss = np.mean((feed_forward - y) ** 2) * 0.5
        return loss

    def calc_layer_errors(self, X, y):
```

```

feed_forward = self.feed_forward(X)

self.layer_errors = []
error_last_layer = (feed_forward - y) * self.activation_grad(self.z_ns[-1])
self.layer_errors.append(error_last_layer)

for i in range(self.num_layers - 2, -1, -1):
    error = self.activation_grad(self.z_ns[i]) * np.dot(self.layer_errors[-1], self.weights[i])
    self.layer_errors.insert(0, error)
return self.layer_errors

def calc_grads(self, X, y):
    self.calc_layer_errors(X, y)
    self.biases_grad, self.weights_grad = [], []

    for i in range(self.num_layers):
        if i == 0:
            wg = np.dot(X.T, self.layer_errors[i]) / len(X)
        else:
            wg = np.dot(self.a_ns[i].T, self.layer_errors[i]) / len(X)
        self.weights_grad.append(wg)

        bg = np.mean(self.layer_errors[i], axis=0, keepdims=True)
        self.biases_grad.append(bg)

def back_prop(self, X, y):
    self.calc_grads(X, y)
    for i in range(self.num_layers):
        self.biases[i] -= self.lr * self.biases_grad[i]
        self.weights[i] -= self.lr * self.weights_grad[i]

```

(a)

```

In [ ]: nn = NN([6, 2, 2], 0.1, tanh, tanh_grad)
# print the initialized weights
print(nn.weights)
print("\nBiases:", nn.biases)

```

```

[array([[0.5488135 , 0.71518937],
        [0.60276338, 0.54488318],
        [0.4236548 , 0.64589411],
        [0.43758721, 0.891773  ],
        [0.96366276, 0.38344152],
        [0.79172504, 0.52889492]]), array([[0.07103606, 0.0871293 ],
        [0.0202184 , 0.83261985]])]

```

```

Biases: [array([[0.56804456, 0.92559664]]), array([[0.77815675, 0.87001215]])]

```

(b)

For debugging: if you use `np.random.random` to initialize weights and biases and set `np.random.seed(0)` . You will get `[[0.64027023 0.25119984]]` as the prediction

```
In [ ]: # print the prediction
input = np.array([[-1, 1, -1, -1, 1, -1]])
res = nn.feed_forward(input)
print(res)
```

```
[[0.64027023 0.25119984]]
```

(c)

Hint: For the definition of layer errors, you can refer to the material in bCourses: [Files > Homework > HW3 > Guide_ANN.pdf](#) . You will use these errors in back-propagation.

For debugging: `[array([[0.08505865, 0.23282824]]), array([[0.48392403, 0.58612371]])]` are errors for two layers.

```
In [ ]: # print layer errors
layer_outputs = []
observed = np.array([-1, -1])
print(nn.calc_layer_errors(input, observed))
```

```
[array([[0.08505865, 0.23282824]]), array([[0.48392403, 0.58612371]])]
```

(d)

```
In [ ]: # print updated weights & biases
nn.back_prop(input, observed)
print("Updated Weights:")
for i, weights in enumerate(nn.weights):
    print(f"Layer {i+1}:")
    print(weights)
print("\nUpdated Biases:")
for i, biases in enumerate(nn.biases):
    print(f"Layer {i+2}:")
    print(biases)
```

Updated Weights:

Layer 1:

```
[[0.55731937 0.73847219]
 [0.59425751 0.52160036]
 [0.43216066 0.66917694]
 [0.44609308 0.91505582]
 [0.9551569 0.3601587 ]
 [0.8002309 0.55217774]]
```

Layer 2:

```
[[0.07428843 0.09106854]
 [0.05552456 0.87538229]]
```

Updated Biases:

Layer 2:

```
[[0.5595387 0.90231381]]
```

Layer 3:

```
[[0.72976435 0.81139978]]
```

Question 2

If you want, you can read the following material to help you with the code!

Let's denote a as the output of our simple perceptron model.

$$\mathbf{a} = \sigma(\mathbf{z}) = \sigma(\mathbf{X}\mathbf{w} + \mathbf{b})$$

where $\mathbf{X} \in \mathbb{R}^{N_{bz} \times N_f}$ is the input data, N_{bz} is the batch size (i.e. number of samples in one calculation) and N_f is the dimension of input features. $\mathbf{w} \in \mathbb{R}^{N_f \times N_o}$ is the weights matrix and $\mathbf{b} \in \mathbb{R}^{N_o}$ is the biases. N_o is the dimension of outputs. σ is the activation function.

The loss function is defined as half of the mean squared error (MSE) between predicted values \mathbf{a} and the ground truth \mathbf{O} :

$$C = \frac{1}{2} \text{MSE}(\mathbf{a}, \mathbf{O}) = \frac{1}{2N_{bz}N_o} \sum_m \sum_i (a_{mi} - O_{mi})^2$$

Using the chain rule to calculate the gradient of weights and biases. First define

$$\delta = \frac{\partial C}{\partial \mathbf{z}} = \frac{\partial C}{\partial \mathbf{a}} \odot \sigma'(\mathbf{z}) = \frac{1}{N_{bz}N_o} (\mathbf{a} - \mathbf{O}) \odot \sigma'(\mathbf{z})$$

where \odot is elementwise product.

$$\frac{\partial C}{\partial \mathbf{w}} = \mathbf{X}^T \frac{\partial C}{\partial \mathbf{z}} = \mathbf{X}^T \delta$$

$$\frac{\partial C}{\partial \mathbf{b}} = \sum_m \delta_m$$

where δ_m is the m -th row of matrix δ .

```
In [ ]: class SimplePerceptron():
        """
        Implementation of simple perceptron regressor
        """
        def __init__(
            self,
            input_dim,
            output_dim,
            learning_rate=0.0001,
            activation=lambda x: x,
            activation_grad=lambda x: 1
        ):
            """
            Initialize a simple perceptron regressor

            Parameters
            -----
            input_dim: int
                Input dimension
            output_dim: int
                Output dimension
            learning_rate: float
                Learning Rate
            activation: Callable
                Activation function
            activation_grad: Callable
                Activation function's first derivative
            """
            self.input_dim = input_dim
            self.output_dim = output_dim
            self.activation = activation
            self.activation_grad = activation_grad
            self.lr = learning_rate
            self.weights = np.random.random((input_dim, output_dim))
            self.biases = np.random.random(output_dim)

        def predict(self, X):
            """
            Get predicted values of given input

            Parameters
            -----
            X: numpy.ndarray
                Input samples, shape (n_samples, input_dim)

            Returns
            -----
            y: numpy.ndarray
                Output, shape (n_samples, output_dim)
            """
            # Convert rank-1 array to rank-2
            if len(X.shape) == 1:
                X = X.reshape((-1, 1))

            # Check that the dimension of accepted input data is the same as expected
```

```

    dim = X.shape[1]
    if dim != self.input_dim:
        raise Exception(f"Expected input size {self.input_dim}, accepted {dim}")

    self.z = np.dot(X, self.weights) + self.biases
    self.a = self.activation(self.z)
    return self.a

def fit(self, X, y):
    """
    Fit the model by doing an steepest descent step

    Parameters
    -----
    X: numpy.ndarray
        Input samples, shape (n_samples, input_dim)
    y: numpy.ndarray
        Labels, shape (n_samples, output_dim)
    """
    self.predict(X)

    errors = (self.a - y) * self.activation_grad(self.z) / y.size #shape (n_s
    weights_grad = np.dot(X.T, errors) #shape (inp
    bias_grad = np.sum(errors, axis=0) #shape (out

    # Update weights and biases from the gradient
    # just do a simple gradient descent
    self.weights -= self.lr * weights_grad
    self.biases -= self.lr * bias_grad

def train_one_epoch(self, X, y, batch_size=32):
    """
    Train the model for one epoch

    Parameters
    -----
    X: numpy.ndarray
        Input samples, shape (n_samples, input_dim)
    y: numpy.ndarray
        Labels, shape (n_samples, output_dim)
    batch_size: int
        Batch size
    """
    if len(X.shape) == 1:
        X = X.reshape((-1, 1))
    if len(y.shape) == 1:
        y = y.reshape((-1, 1))

    indices = np.arange(X.shape[0])
    np.random.shuffle(indices)
    for i in range(0, X.shape[0], batch_size):
        batch = indices[i:i+batch_size]
        self.fit(X[batch], y[batch])

def evaluate(self, X, y):
    """

```

```

    Get the loss function with given data

    Parameters
    -----
    X: numpy.ndarray
        Input samples, shape (n_samples, input_dim)
    y: numpy.ndarray
        Labels, shape (n_samples, output_dim)

    Returns
    -----
    L: float
        Loss function value.
    """
    # Transform the single-sample data into 2-dimensional, for the convenience
    if len(X.shape) == 1:
        X = X.reshape((-1, 1))
    if len(y.shape) == 1:
        y = y.reshape((-1, 1))

    y_pred = self.predict(X)
    # mean squared error
    return np.mean((y_pred - y) ** 2)

def get_weights(self):
    """
    Get weights and biases
    """
    return self.weights.copy(), self.biases.copy()

def set_weights(self, weights, biases):
    """
    Set weights and biases
    """
    self.weights = weights
    self.biases = biases

```

(a)

Hint: Think about what activation function in the SimplePerceptron model should be used to perform a logistic regression?

One-dimension logistic regression expression:

$$y = \frac{1}{1 + e^{-(ax+b)}}$$

```

In [ ]: import pandas as pd
        from sklearn.preprocessing import OneHotEncoder
        data = pd.read_csv('Datasets/titanic.csv')
        data = data.dropna()

        categorical_features = data[['Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch', 'Ticket', '

```

```
survival_status = data[['Survived']]
encoder = OneHotEncoder()
categorical_encoded = encoder.fit_transform(categorical_features).toarray()
survival_encoded = encoder.fit_transform(survival_status).toarray()

print("Categories:\n", categorical_encoded)
print("Survival status:\n", survival_encoded)
```

C:\Users\artui\AppData\Local\Temp\ipykernel_20244\3739233400.py:1: DeprecationWarning:

Pyarrow will become a required dependency of pandas in the next major release of pandas (pandas 3.0),

(to allow more performant data types, such as the Arrow string type, and better interoperability with other libraries)

but was not found to be installed on your system.

If this would cause problems for you,

please provide us feedback at <https://github.com/pandas-dev/pandas/issues/54466>

```
import pandas as pd
```


Categories:

```
[[1. 0. 0. ... 1. 0. 0.]  
[1. 0. 0. ... 0. 0. 1.]  
[1. 0. 0. ... 0. 0. 1.]  
...  
[1. 0. 0. ... 1. 0. 0.]  
[1. 0. 0. ... 0. 0. 1.]  
[1. 0. 0. ... 1. 0. 0.]]
```

Survival status:

```
[[0. 1.]  
[0. 1.]  
[1. 0.]  
[0. 1.]  
[0. 1.]  
[0. 1.]  
[0. 1.]  
[1. 0.]  
[0. 1.]  
[1. 0.]  
[1. 0.]  
[0. 1.]  
[1. 0.]  
[0. 1.]  
[1. 0.]  
[1. 0.]  
[0. 1.]  
[1. 0.]  
[1. 0.]  
[1. 0.]  
[1. 0.]  
[0. 1.]  
[1. 0.]  
[0. 1.]  
[1. 0.]  
[0. 1.]  
[1. 0.]  
[1. 0.]  
[1. 0.]  
[0. 1.]  
[0. 1.]  
[0. 1.]  
[0. 1.]  
[1. 0.]  
[0. 1.]  
[0. 1.]  
[0. 1.]  
[0. 1.]  
[0. 1.]  
[1. 0.]  
[0. 1.]  
[0. 1.]  
[0. 1.]  
[0. 1.]  
[1. 0.]  
[0. 1.]  
[1. 0.]  
[1. 0.]  
[0. 1.]  
[1. 0.]  
[1. 0.]
```

10/20

[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[1. 0.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[1. 0.]
[1. 0.]
[0. 1.]

```
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[1. 0.]
[1. 0.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]
```

(b)

Here, we will do the following things:

- Split the data to k -folds, use the $k - 1$ folds for training and the other fold for testing
- In each fold, further split the training data ($k - 1$ folds of the whole dataset) by 80%/20%. Use the 80% to train the model in certain epochs and use the 20% as a validation set to determine the best model during training.
- For each fold, plot the validation loss against training
- For each fold, report the mean squared error of the total training set (the $k - 1$ fold data) and the test set (the other fold data)

```
In [ ]: from sklearn.model_selection import train_test_split, KFold
import matplotlib.pyplot as plt
import numpy as np

def KFoldCrossValidation(
    k, X, y, epochs, draw_curve=True, learning_rate=0.0001,
    activation=lambda x: x, activation_grad=lambda x: 1,
    verbose=True
):
    """
    K-Fold Validation

    Parameters
    -----
    k: int
```

```

    The number of fold in validation
X: numpy.ndarray
    Input samples, shape (n_samples, input_dim)
y: numpy.ndarray
    Labels, shape (n_samples, output_dim)
epochs: int
    Number of training epochs
draw_curve: bool
    Whether to draw the validation loss against training steps
learning_rate: float
    Learning rate
activation: Callable
    Activation function
activation_grad: Callable
    Gradient of activate function
verbose: bool
    Whether to print information

Returns
-----
model: list
    List of k models
"""
# Reshape the X,y if they are rank-1
if len(X.shape) == 1:
    X = X.reshape((-1, 1))
if len(y.shape) == 1:
    y = y.reshape((-1, 1))
input_dim = X.shape[1]
output_dim = y.shape[1]

# do a 5-fold validation
kf = KFold(n_splits=k)

if draw_curve:
    n_row, n_col = int(np.ceil(k / 3)), 3
    fig, axes = plt.subplots(n_row, n_col, figsize=(5 * n_col, 5 * n_row), cons
    axes = axes.flatten()

models = []
train_err_list, test_err_list = [], []
for i, (train_selector, val_selector) in enumerate(kf.split(X)):
    # Decide training examples and testing examples for this fold
    X_train, X_test = X[train_selector], X[val_selector]
    y_train, y_test = y[train_selector], y[val_selector]

    # define you models here
    model = SimplePerceptron(input_dim, output_dim, learning_rate, activation,

    # futher split the X_train by 80%/20%
    # use train_test_split function in sklearn
    X_train_in, X_val, y_train_in, y_val = train_test_split(X_train, y_train, t

    lowest_val_err = np.inf
    val_err_list = []
    for _ in range(epochs):

```

```

        # Train model on a number of epochs, and test performance in the valida
        model.train_one_epoch(X_train_in, y_train_in)
        val_err = model.evaluate(X_val, y_val)
        val_err_list.append(val_err)
        if val_err < lowest_val_err:
            lowest_val_err = val_err
            weights, biases = model.get_weights()

    # The final number of epochs is when the minimum error in validation set oc
    if verbose:
        print("Number of epochs with lowest validation:", np.argmin(val_err_list))
        model.set_weights(weights, biases)
        models.append(model)

    # Report MSE on X_train and X_test
    train_err = model.evaluate(X_train, y_train)
    train_err_list.append(train_err)
    test_err = model.evaluate(X_test, y_test)
    test_err_list.append(test_err)

    if draw_curve:
        axes[i].plot(np.arange(epochs), np.log10(val_err_list), label='Validati
        axes[i].set_xlabel('Epochs')
        axes[i].set_ylabel('Log Loss')
        axes[i].legend()
        axes[i].set_title(f"Epoch {i}")

    if verbose:
        print("Final results:")
        print(f"Training error: {np.mean(train_err_list)}+/-{np.std(train_err_list)}")
        print(f"Testing error: {np.mean(test_err_list)}+/-{np.std(test_err_list)}")

    return models

```

```

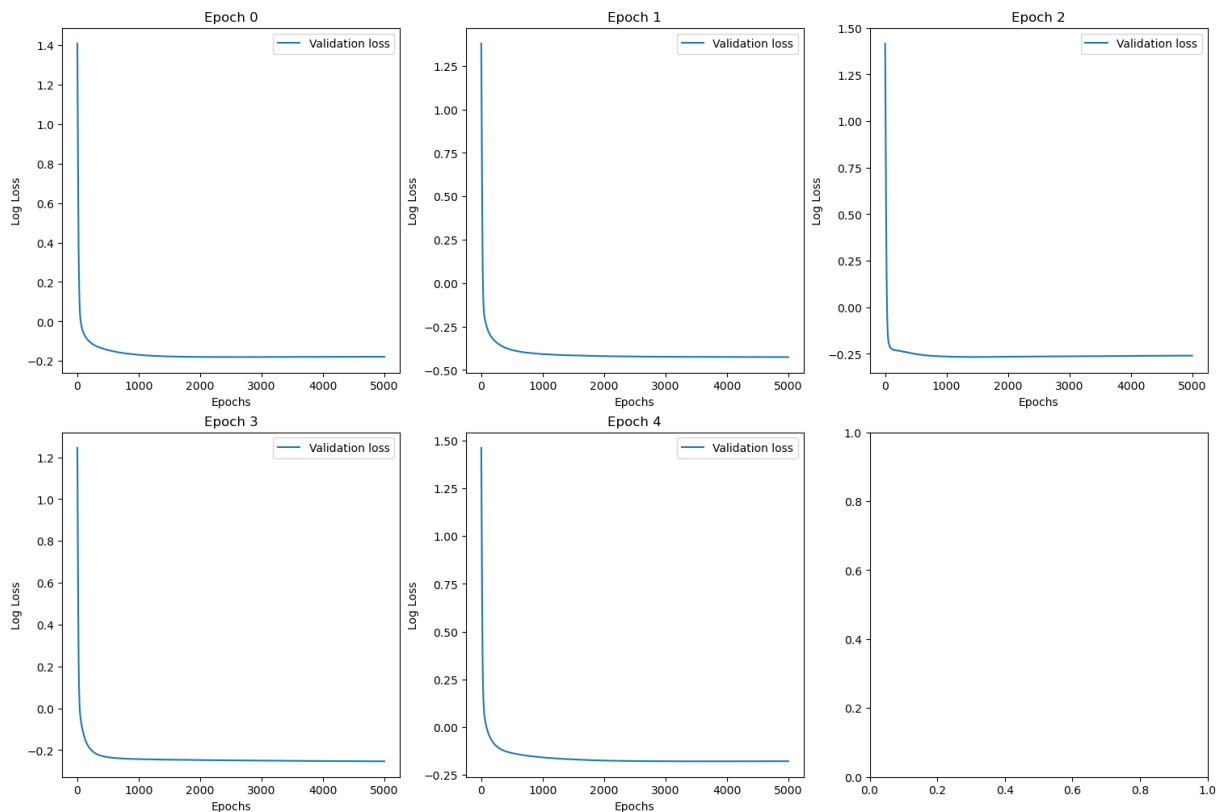
In [ ]: result = KFoldCrossValidation(k=5, X=categorical_encoded, y=survival_encoded, learn
# Increased Learning rate by factor of 100 b/c # of epochs was too large.

```

```

Number of epochs with lowest validation: 2582
Number of epochs with lowest validation: 4982
Number of epochs with lowest validation: 1443
Number of epochs with lowest validation: 4998
Number of epochs with lowest validation: 3487
Final results:
Training error: 0.12778737674083654+/-0.02770040001834124
Testing error: 0.5494929936374364+/-0.05799171569291407

```



Question 3

(a)

Note: `KFoldCrossValidation` returns a model list, just pick any one of them (for example, the first) for prediction on test dataset.

```
In [ ]: def generate_X(number, boundary):
        x = (np.random.random(number) * 2 - 1) * boundary
        return x

        def generate_data(number, stochasticity=0.2):
            x = generate_X(number, 10)
            ratio = generate_X(number, stochasticity) + 1
            y = 3 * np.sin(x) * ratio + 5
            return x, y
```

```
In [ ]: def show_correlation(y_pred, y_true):
        y_pred = y_pred.flatten()
        y_true = y_true.flatten()
        r = [
            np.min([np.min(y_pred), np.min(y_true)]),
            np.max([np.max(y_pred), np.max(y_true)])
        ]
        fig, ax = plt.subplots(1, 1, figsize=(5, 4))
        plt.scatter(y_pred, y_true, s=5)
        ax.plot(r, r, color='red')
```

```

ax.set_xlabel("Predictions")
ax.set_ylabel("Ground truth")
corr = np.corrcoef([y_pred, y_true])[1, 0]
print("Correlation coefficient:", corr)

def show_plot(x, y_true, y_pred=None):
    fig, ax = plt.subplots(1, 1, figsize=(5, 4))
    ax.scatter(x, y_true, s=0.5, label='Ground Truth')
    if y_pred is not None:
        ax.scatter(x, y_pred, s=0.5, label="Predicted")
    ax.legend()

```

```

In [ ]: X, y = generate_data(5000)
        kfold_res = KFoldCrossValidation(5, X, y, 2000)

```

Number of epochs with lowest validation: 434

Number of epochs with lowest validation: 472

Number of epochs with lowest validation: 400

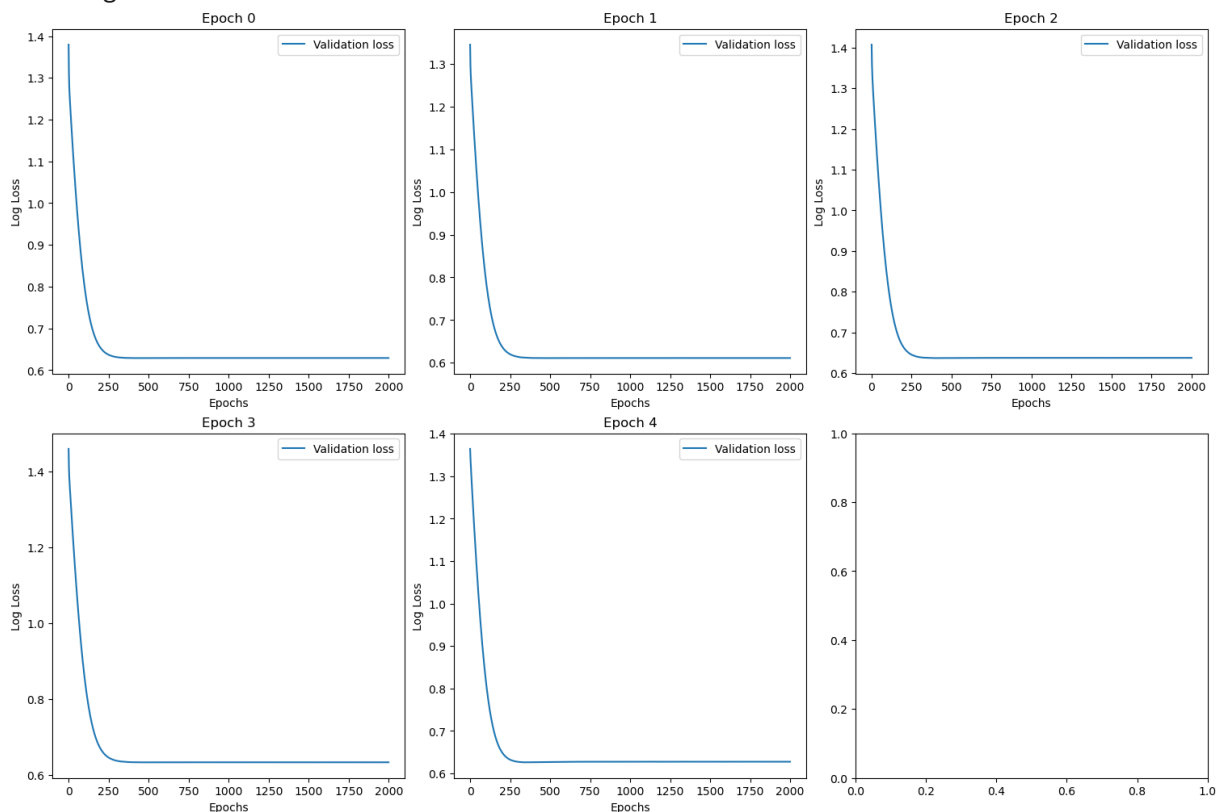
Number of epochs with lowest validation: 501

Number of epochs with lowest validation: 354

Final results:

Training error: 4.272383649695879+/-0.028393052413088742


Testing error: 4.274331762526268+/-0.11630963716861877



Q. How well is the prediction? A. The correlation coefficient is low (nearly 0), indicating that there is almost no correlation. There is also significant variance between the line of regression and the predictions in blue. This does not fit the prediction well. Our simple perceptron cannot learn non-linear (i.e., random) data.


```
In [ ]: X_test, y_test = generate_data(1000)
        show_correlation(kfold_res[0].predict(X_test), y_test)
```

Correlation coefficient: 0.13966518520677262



(b)

```
In [ ]: from sklearn.neural_network import MLPRegressor
        from sklearn.metrics import mean_squared_error as mse

        def KFoldCrossValidationMLP(
            k, X, y, epochs,
            learning_rate=0.01,
            activation='tanh',
            hidden_layers=(8,),
            verbose=True
        ):
            """
            K-Fold Validation for MLPRegressor

            Parameters
            -----
            k: int
                The number of fold in validation
            X: numpy.ndarray
                Input samples, shape (n_samples, input_dim)
            y: numpy.ndarray
                Labels, shape (n_samples, output_dim)
            epochs: int
                Number of training epochs
            kwargs: keyword arguments
                Arguments to init a SimplePerceptron model. i.e. learning_rate, activation,
```

```

Returns
-----
model: SimplePerceptron
    The best model
"""
# MLPRegressor takes 2D array for features
# and 1D array for labels if the dimension of output is 1
if len(X.shape) == 1:
    X = X.reshape(-1, 1)

# do a 5-fold validation
kf = KFold(n_splits=k, shuffle=True)

models = []
train_err_list, test_err_list = [], []
for i, (train_selector, val_selector) in enumerate(kf.split(np.arange(X.shape[0]
    # Decide training examples and testing examples for this fold
    X_train, X_test = X[train_selector], X[val_selector]
    y_train, y_test = y[train_selector], y[val_selector]

    model = MLPRegressor(
        max_iter=epochs,
        activation=activation,
        learning_rate_init=learning_rate,
        hidden_layer_sizes=hidden_layers,
        early_stopping=True, validation_fraction=0.2, learning_rate='constant',
    )
    model.fit(X_train, y_train)
    models.append(model)

    # Report MSE on X_train and X_test
    y_train_pred = model.predict(X_train)
    train_err = mse(y_train_pred, y_train)
    train_err_list.append(train_err)

    y_test_pred = model.predict(X_test)
    test_err = mse(y_test_pred, y_test)
    test_err_list.append(test_err)
    if verbose:
        print(f"Train error: {train_err}, Test error: {test_err}")

if verbose:
    print("Final results:")
    print(f"Training error: {np.mean(train_err_list)}+/-{np.std(train_err_list)}")
    print(f"Testing error: {np.mean(test_err_list)}+/-{np.std(test_err_list)}")

return models[np.argmin(test_err_list)]


```

```

In [ ]: # X and y are the same from the simple perceptron, so that we can visualize the dif
mlp = KFoldCrossValidationMLP(5, X, y, epochs=2000)
show_correlation(mlp.predict(X_test.reshape(-1, 1)), y_test)

```

Train error: 2.3500982121194642, Test error: 2.4767437507469303
 Train error: 0.34262626618828385, Test error: 0.44308386463217914
 Train error: 2.3473432906148637, Test error: 2.3682804199953615
 Train error: 2.365629861964926, Test error: 2.2562631545457874
 Train error: 0.37080745551727784, Test error: 0.33845115030673945
 Final results:
 Training error: 1.5553010172809631+/-0.9787003044850086
 Testing error: 1.5765644680453996+/-0.9712703083985693
 Correlation coefficient: 0.9568416424156395




The correlation is much, much better. The coefficient is nearly 1, which would be perfect. The problem with the simple perceptron is that we applied a simple sigmoid function, when MLPs generally use RELU functions. In addition, we apply multiple neurons for our single-layer MLP, which helps inform our data better.

(c)

```
In [ ]: mlp = KFoldCrossValidationMLP(5, X, y, hidden_layers=(8, 4), epochs=2000)
        show_correlation(mlp.predict(X_test.reshape(-1, 1)), y_test)
```

Train error: 0.06403303809927906, Test error: 0.06857576685710352
 Train error: 0.07649579626310962, Test error: 0.07881612955313179
 Train error: 0.07100518849340276, Test error: 0.07279177989859721
 Train error: 0.06258035293811803, Test error: 0.06206156944405865
 Train error: 0.07067734435754343, Test error: 0.06736115953903919
 Final results:
 Training error: 0.06895834403029058+/-0.005077304982551014
 Testing error: 0.06992128105838608+/-0.005612446715225064
 Correlation coefficient: 0.9932367451919568



A good MLP usually has 2-3 layers. Here, I give it a decent amount of neurons per layer, with 2 layers. I was able to achieve a better correlation coefficient.